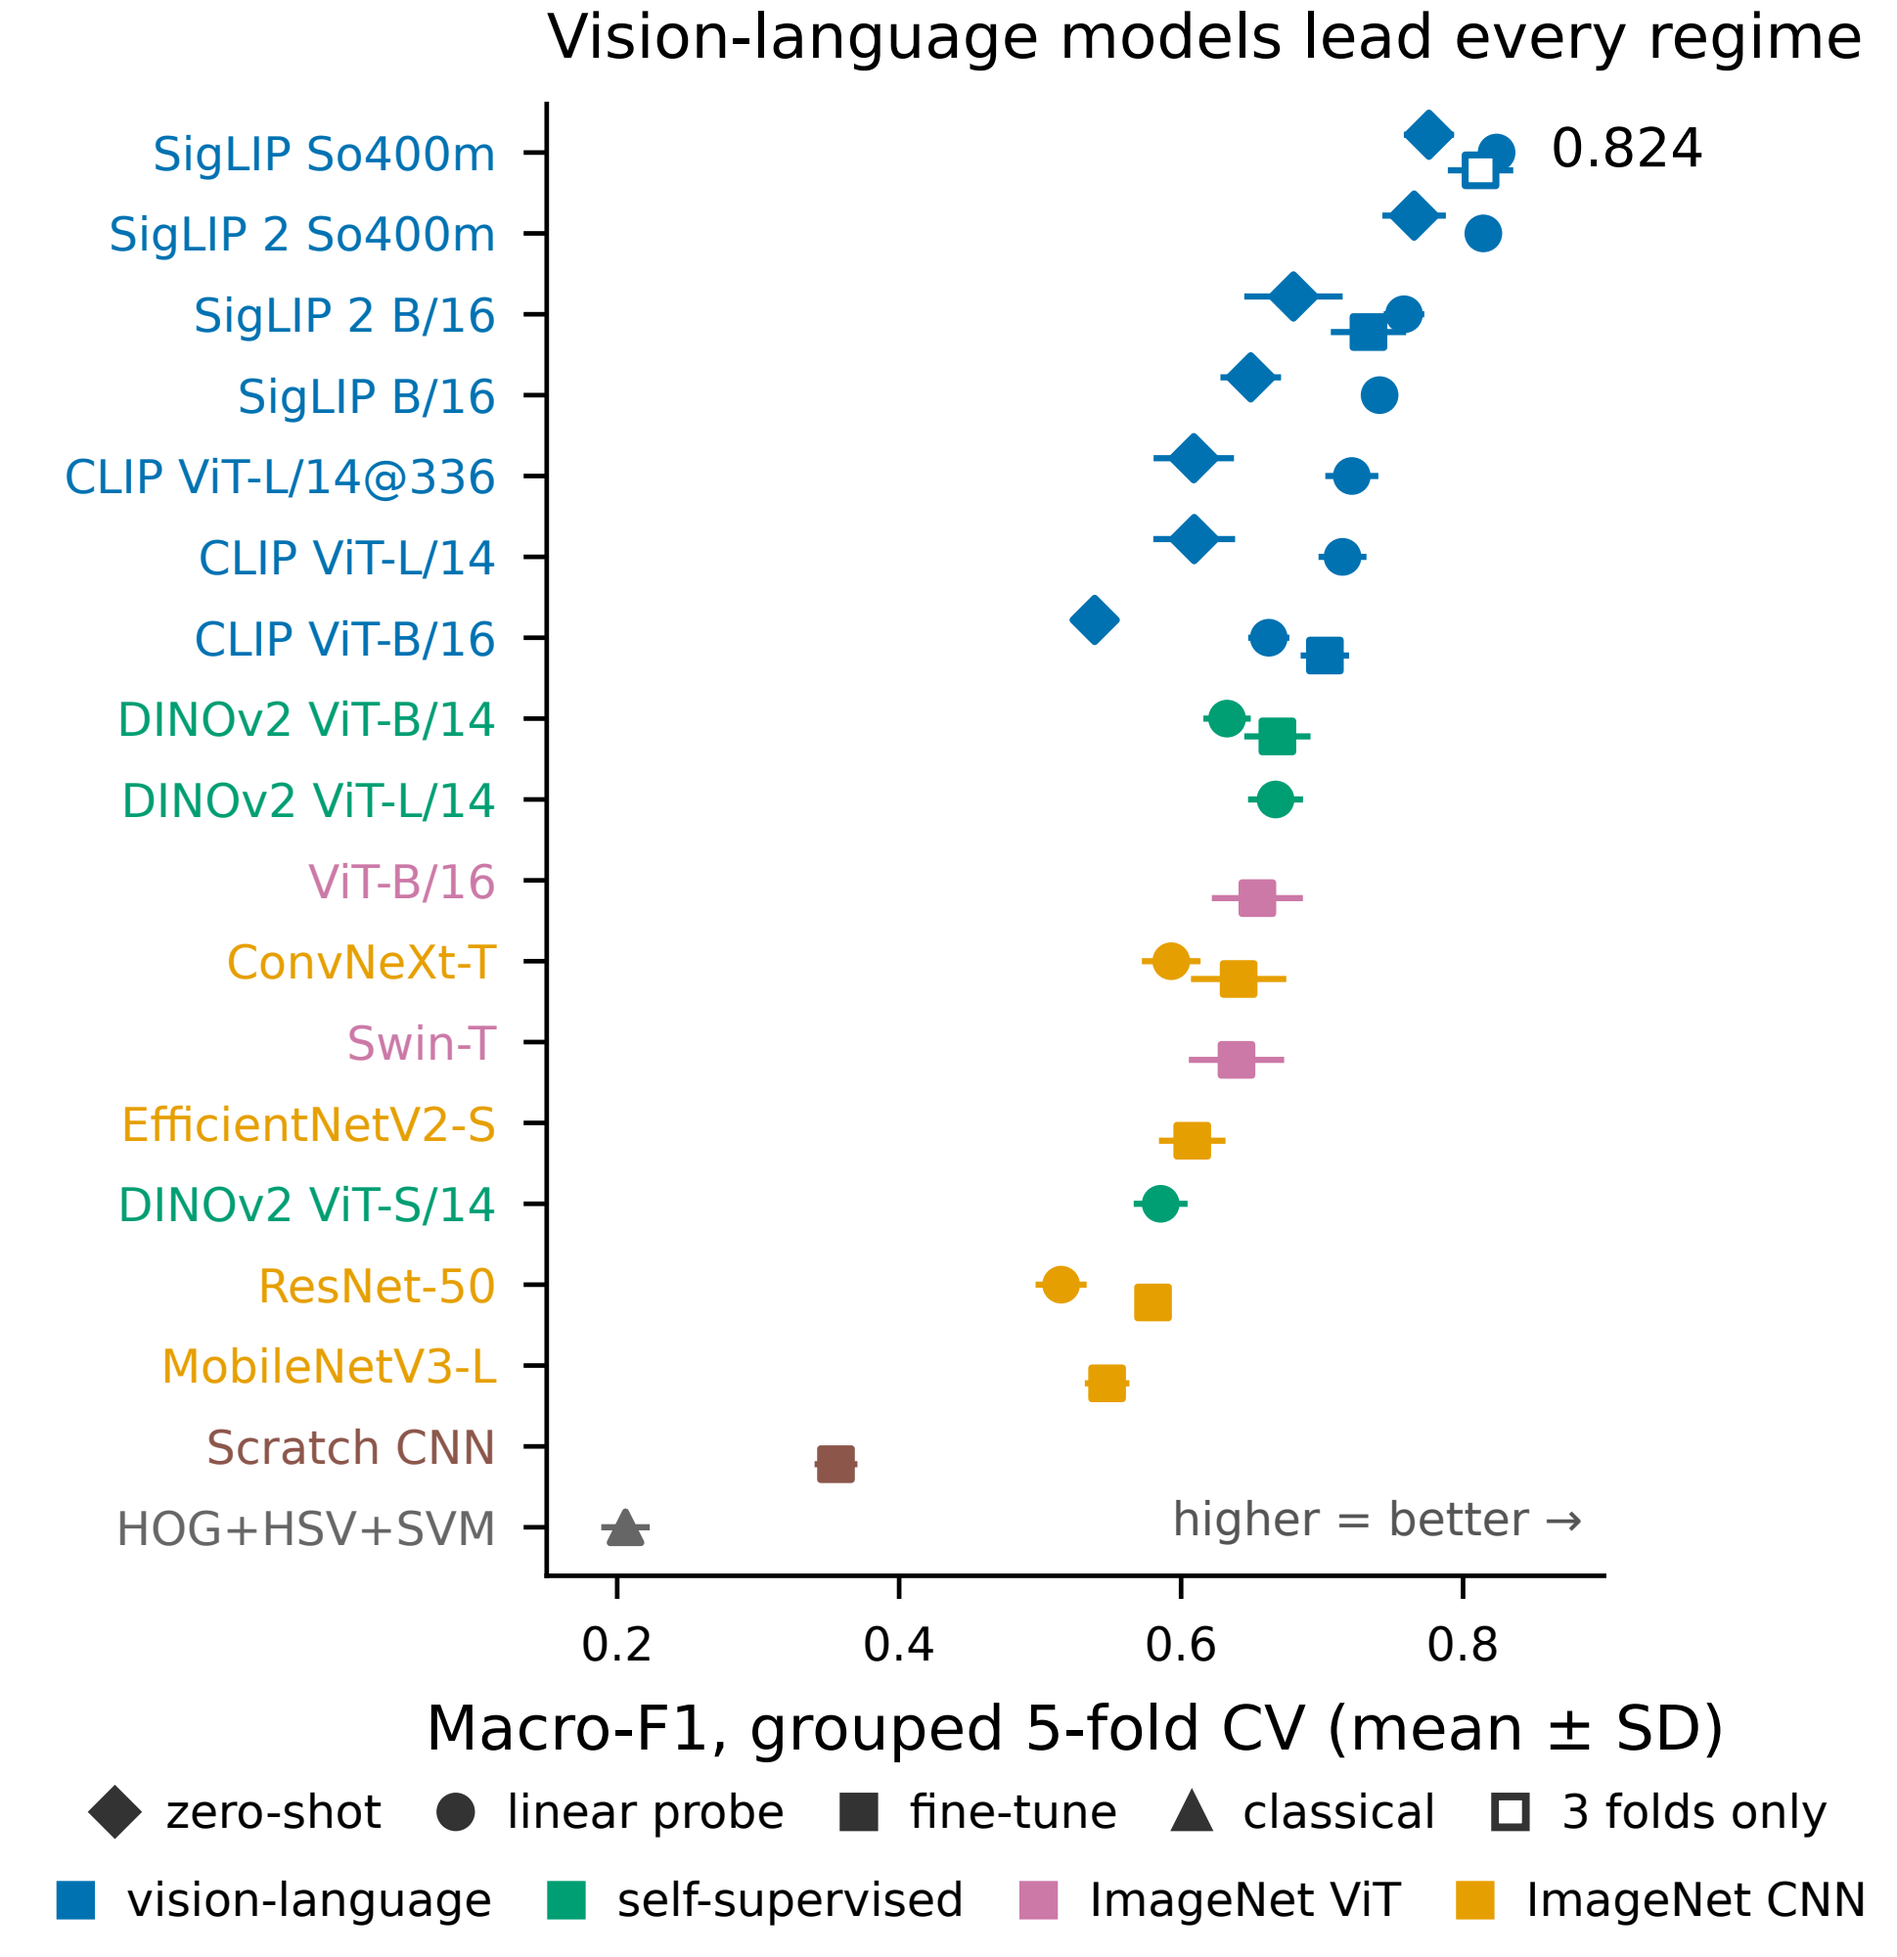
Vision-language models lead every regime
SigLIP So400m
SigLIP 2 So400m
SigLIP 2 B/16
SigLIP B/16
CLIP ViT-L/14@336
CLIP ViT-L/14
CLIP ViT-B/16
DINOv2 ViT-B/14
DINOv2 ViT-L/14
ViT-B/16
ConvNeXt-T
Swin-T
EfficientNetV2-S
DINOv2 ViT-S/14
ResNet-50
MobileNetV3-L
Scratch CNN
HOG+HSV+SVM
0.824
higher = better →
0.2
0.4
0.6
0.8
Macro-F1, grouped 5-fold CV (mean ± SD)
zero-shot
linear probe
fine-tune
classical
3 folds only
vision-language
self-supervised
ImageNet ViT
ImageNet CNN