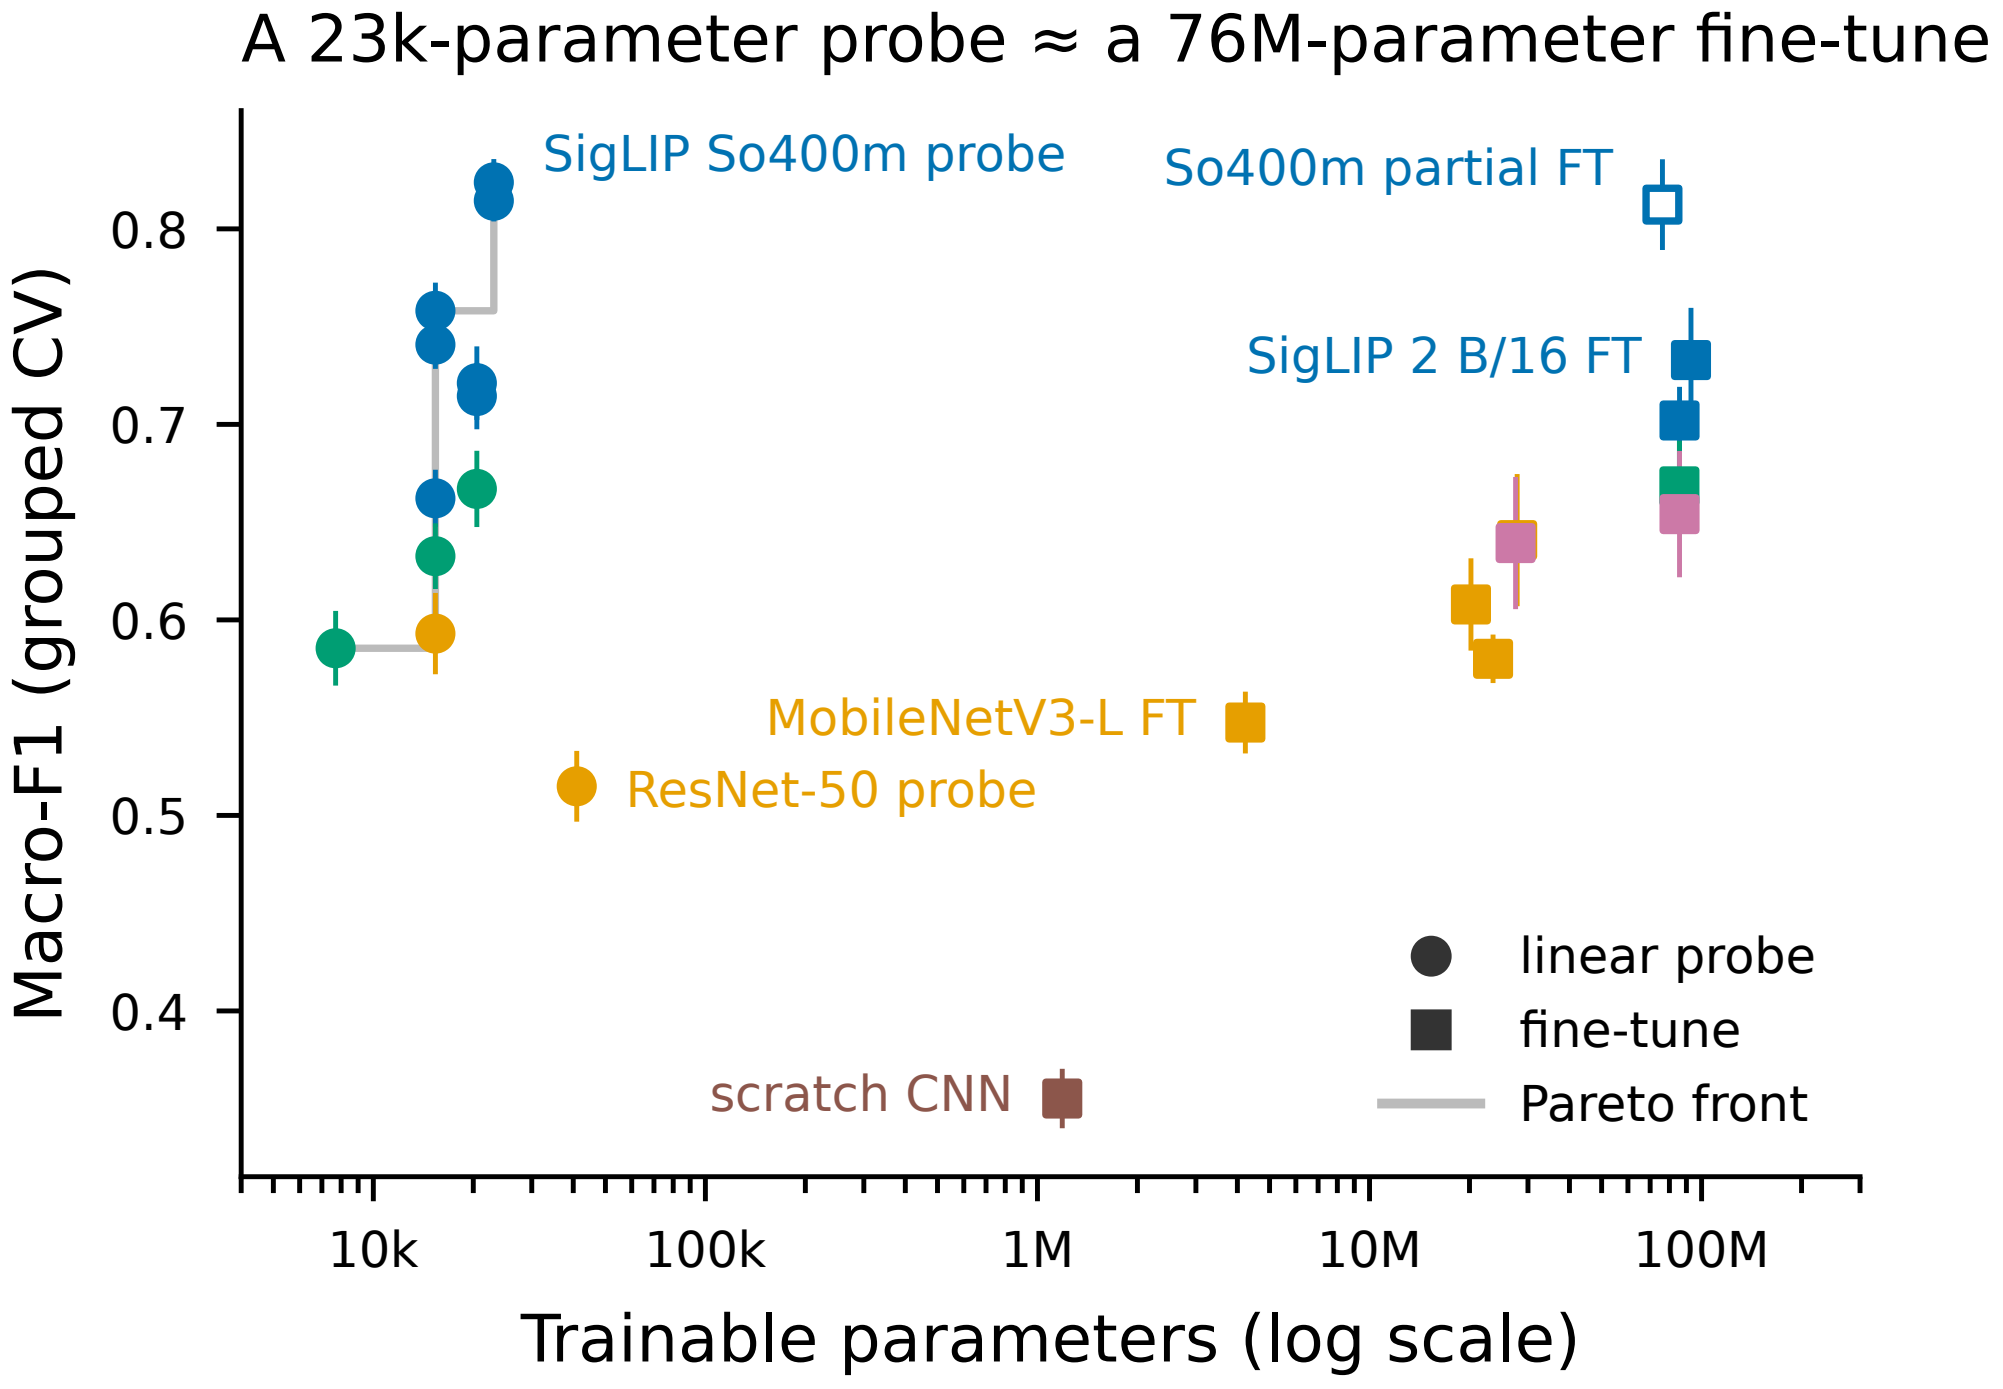
A 23k-parameter probe ≈ a 76M-parameter fine-tune
SigLIP So400m probe
So400m partial FT
SigLIP 2 B/16 FT
MobileNetV3-L FT
ResNet-50 probe
scratch CNN
linear probe
fine-tune
Pareto front
Macro-F1 (grouped CV)
0.8
0.7
0.6
0.5
0.4
10k
100k
1M
10M
100M
Trainable parameters (log scale)