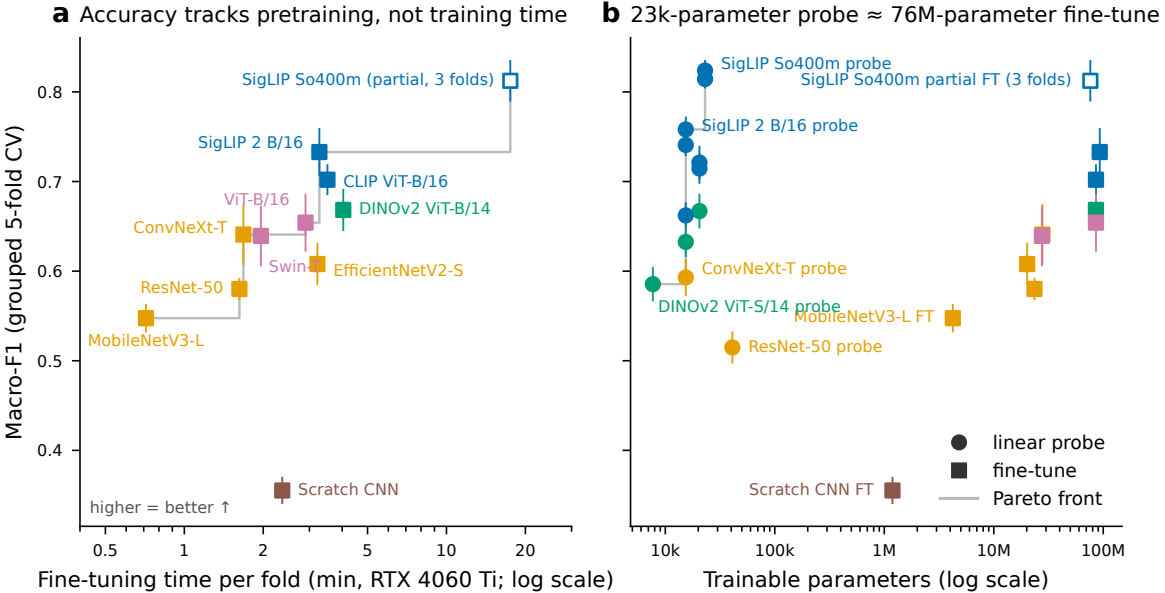
a Accuracy tracks pretraining, not training time
b 23k-parameter probe ≈ 76M-parameter fine-tune
Macro-F1 (grouped 5-fold CV)
0.8
0.7
0.6
0.5
0.4
SigLIP So400m (partial, 3 folds)
SigLIP 2 B/16
CLIP ViT-B/16
DINOv2 ViT-B/14
ViT-B/16
ConvNeXt-T
Swin-T
EfficientNetV2-S
ResNet-50
MobileNetV3-L
Scratch CNN
higher = better ↑
0.5
1
2
5
10
20
Fine-tuning time per fold (min, RTX 4060 Ti; log scale)
SigLIP So400m probe
SigLIP So400m partial FT (3 folds)
SigLIP 2 B/16 probe
ConvNeXt-T probe
DINOv2 ViT-S/14 probe
MobileNetV3-L FT
ResNet-50 probe
Scratch CNN FT
linear probe
fine-tune
Pareto front
10k
100k
1M
10M
100M
Trainable parameters (log scale)